

Untitled0

December 4, 2022

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[1]: import tensorflow as tf
import numpy as np

batch_size = 32
num_classes = 3
epochs = 20

train_ds = tf.keras.utils.image_dataset_from_directory(
    directory='/content/drive/MyDrive/ML/DataSetFruits/',
    validation_split=0.2,
    subset="training",
    labels='inferred',
    label_mode='categorical',
    seed=123,
    shuffle=True,
    batch_size=batch_size,
    image_size=(100,100))


val_ds = tf.keras.utils.image_dataset_from_directory(
    directory='/content/drive/MyDrive/ML/DataSetFruits/',
    validation_split=0.2,
    subset="validation",
    labels='inferred',
    label_mode='categorical',
    seed=123,
    shuffle=True,
    batch_size=batch_size,
    image_size=(100,100))
```

Found 1472 files belonging to 3 classes.
Using 1178 files for training.
Found 1472 files belonging to 3 classes.

Using 294 files for validation.

```
[3]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.axis("off")
```



The goal is to differentiate between apples, bananas, and cocos

1 Basic Sequential Model

```
[190]: model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(100, 100, 3)),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(num_classes, activation='softmax'),
])
```

Let's check the summary

```
[191]: model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
flatten_7 (Flatten)	(None, 30000)	0
dense_21 (Dense)	(None, 512)	15360512
dropout_20 (Dropout)	(None, 512)	0
dense_22 (Dense)	(None, 512)	262656
dropout_21 (Dropout)	(None, 512)	0
dense_23 (Dense)	(None, 3)	1539

=====
Total params: 15,624,707
Trainable params: 15,624,707
Non-trainable params: 0
=====

Now we compile and then fit to our dataset

```
[189]: model.compile(loss='categorical_crossentropy',
    optimizer='rmsprop',
    metrics=['accuracy'])

history = model.fit(train_ds,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=val_ds)
```

Epoch 1/20

30/37 [=====>...] - ETA: 1s - loss: 4634.1978 - accuracy:
0.6240

```

↳ -----
KeyboardInterrupt                                Traceback (most recent call↳
↳ last)

<ipython-input-189-ccdc65fc173d> in <module>
      3             metrics=['accuracy'])
      4
----> 5 history = model.fit(train_ds,
      6                     batch_size=batch_size,
      7                     epochs=epochs,

/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.py in ↳
↳ error_handler(*args, **kwargs)
      62         filtered_tb = None
      63         try:
---> 64         return fn(*args, **kwargs)
      65     except Exception as e: # pylint: disable=broad-except
      66         filtered_tb = _process_traceback_frames(e.__traceback__)

/usr/local/lib/python3.8/dist-packages/keras/engine/training.py in ↳
↳ fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, ↳
↳ validation_data, shuffle, class_weight, sample_weight, initial_epoch, ↳
↳ steps_per_epoch, validation_steps, validation_batch_size, validation_freq, ↳
↳ max_queue_size, workers, use_multiprocessing)
    1407         _r=1):
    1408             callbacks.on_train_batch_begin(step)
-> 1409             tmp_logs = self.train_function(iterator)
    1410             if data_handler.should_sync:
    1411                 context.async_wait()

/usr/local/lib/python3.8/dist-packages/tensorflow/python/util/
↳ traceback_utils.py in error_handler(*args, **kwargs)
    148         filtered_tb = None
    149         try:
--> 150         return fn(*args, **kwargs)
    151     except Exception as e:
    152         filtered_tb = _process_traceback_frames(e.__traceback__)
```

```

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/
↳def_function.py in __call__(self, *args, **kwargs)
    913
    914         with OptionalXlaContext(self._jit_compile):
--> 915             result = self._call(*args, **kwargs)
    916
    917             new_tracing_count = self.experimental_get_tracing_count()

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/
↳def_function.py in _call(self, *args, **kwargs)
    945         # In this case we have created variables on the first call, so
↳we run the
    946         # defunned version which is guaranteed to never create
↳variables.
--> 947         return self._stateless_fn(*args, **kwargs) # pylint:
↳disable=not-callable
    948         elif self._stateful_fn is not None:
    949             # Release the lock early so that multiple threads can perform
↳the call

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/function.
↳py in __call__(self, *args, **kwargs)
    2451         (graph_function,
    2452          filtered_flat_args) = self._maybe_define_function(args,
↳kwargs)
-> 2453         return graph_function._call_flat(
    2454             filtered_flat_args, captured_inputs=graph_function.
↳captured_inputs) # pylint: disable=protected-access
    2455

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/function.
↳py in _call_flat(self, args, captured_inputs, cancellation_manager)
    1858         and executing_eagerly):
    1859         # No tape is watching; skip to running the function.
-> 1860         return self._build_call_outputs(self._inference_function.call(
    1861             ctx, args, cancellation_manager=cancellation_manager))
    1862         forward_backward = self._select_forward_and_backward_functions(

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/function.
↳py in call(self, ctx, args, cancellation_manager)
    495         with _InterpolateFunctionError(self):
    496             if cancellation_manager is None:

```

```

--> 497         outputs = execute.execute(
498             str(self.signature.name),
499             num_outputs=self._num_outputs,

/usr/local/lib/python3.8/dist-packages/tensorflow/python/eager/execute.
py in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    52     try:
    53         ctx.ensure_initialized()
--> 54         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name,
py op_name,
    55                                     inputs, attrs, num_outputs)
    56     except core._NotOkStatusException as e:

```

KeyboardInterrupt:

Let's go ahead and graph it

```
[ ]: history.history.keys()
```

```
[ ]: import matplotlib.pyplot as plt

# Plot training & validation accuracy values
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```

```
[ ]: score = model.evaluate(val_ds, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

2 CNN

```
[48]: num_filters = 8
filter_size = 3
pool_size = 2

model = tf.keras.models.Sequential(
    [
        tf.keras.Input(shape=(100, 100, 3)),

```

```

        tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
        tf.keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(num_classes, activation="softmax"),
    ]
)

```

[49]: `model.summary()`

Model: "sequential_7"

Layer (type)	Output Shape	Param #

conv2d_2 (Conv2D)	(None, 98, 98, 32)	896
max_pooling2d_2 (MaxPooling 2D)	(None, 49, 49, 32)	0
conv2d_3 (Conv2D)	(None, 47, 47, 64)	18496
max_pooling2d_3 (MaxPooling 2D)	(None, 23, 23, 64)	0
flatten_5 (Flatten)	(None, 33856)	0
dropout_9 (Dropout)	(None, 33856)	0
dense_13 (Dense)	(None, 3)	101571
=====		
Total params: 120,963		
Trainable params: 120,963		
Non-trainable params: 0		

[50]: `model.compile(loss='categorical_crossentropy',
optimizer='adam',
metrics=['accuracy'])`

`history = model.fit(train_ds,
batch_size=batch_size,
epochs=epochs,
verbose=1,
validation_data=val_ds)`

Epoch 1/20
37/37 [=====] - 21s 552ms/step - loss: 42.4354 -
accuracy: 0.8065 - val_loss: 1.1198e-05 - val_accuracy: 1.0000

Epoch 2/20
37/37 [=====] - 18s 488ms/step - loss: 6.4997e-04 -
accuracy: 0.9992 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 3/20
37/37 [=====] - 18s 487ms/step - loss: 2.0239e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 4/20
37/37 [=====] - 18s 488ms/step - loss: 9.1077e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 5/20
37/37 [=====] - 18s 490ms/step - loss: 1.0120e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 6/20
37/37 [=====] - 18s 489ms/step - loss: 2.0239e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 7/20
37/37 [=====] - 18s 486ms/step - loss: 2.8335e-09 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 8/20
37/37 [=====] - 18s 487ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 9/20
37/37 [=====] - 18s 490ms/step - loss: 3.0359e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 10/20
37/37 [=====] - 19s 506ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 11/20
37/37 [=====] - 18s 484ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 12/20
37/37 [=====] - 18s 486ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 13/20
37/37 [=====] - 18s 485ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 14/20
37/37 [=====] - 18s 484ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 15/20
37/37 [=====] - 18s 488ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

Epoch 16/20
37/37 [=====] - 18s 485ms/step - loss: 3.0359e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000


```


Epoch 17/20
37/37 [=====] - 18s 482ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000
Epoch 18/20
37/37 [=====] - 18s 486ms/step - loss: 0.0000e+00 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000
Epoch 19/20
37/37 [=====] - 18s 485ms/step - loss: 4.0479e-10 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000
Epoch 20/20
37/37 [=====] - 20s 544ms/step - loss: 1.1941e-08 -
accuracy: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000

```

```

[51]: # Plot training & validation accuracy values
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



```
[52]: score = model.evaluate(val_ds, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Test loss: 0.0
Test accuracy: 1.0


3 Pretrained Model and Transfer Learning

```
[4]: train_ds = tf.keras.utils.image_dataset_from_directory(
    directory='/content/drive/MyDrive/ML/DataSetFruits/',
    validation_split=0.2,
    subset="training",
    labels='inferred',
    label_mode='categorical',
    seed=123,
    shuffle=True,
    batch_size=batch_size,
    image_size=(100,100))

val_ds = tf.keras.utils.image_dataset_from_directory(
    directory='/content/drive/MyDrive/ML/DataSetFruits/',
    validation_split=0.2,
    subset="validation",
    labels='inferred',
    label_mode='categorical',
    seed=123,
    shuffle=True,
    batch_size=batch_size,
    image_size=(100,100))

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.axis("off")
```

Found 1472 files belonging to 3 classes.
Using 1178 files for training.
Found 1472 files belonging to 3 classes.
Using 294 files for validation.



```
[5]: val_batches = tf.data.experimental.cardinality(val_ds)
test_ds = val_ds.take(val_batches // 5)
val_ds = val_ds.skip(val_batches // 5)
```

```
[6]: print('Number of validation batches: %d' % tf.data.experimental.
      ↪cardinality(val_ds))
print('Number of test batches: %d' % tf.data.experimental.cardinality(test_ds))
```

```
Number of validation batches: 8
Number of test batches: 2
```

```
[7]: AUTOTUNE = tf.data.AUTOTUNE
```

```
train_dataset = train_ds.prefetch(buffer_size=AUTOTUNE)
validation_dataset = val_ds.prefetch(buffer_size=AUTOTUNE)
test_dataset = test_ds.prefetch(buffer_size=AUTOTUNE)
```

```
[8]: data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip('horizontal'),
    tf.keras.layers.RandomRotation(0.2),
])
```

```
[9]: for image, _ in train_dataset.take(1):
    plt.figure(figsize=(10, 10))
    first_image = image[0]
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        augmented_image = data_augmentation(tf.expand_dims(first_image, 0))
        plt.imshow(augmented_image[0] / 255)
        plt.axis('off')
```



```
[10]: tf.keras.applications.mobilenet_v2.MobileNetV2(  
    input_shape=None,  
    alpha=1.0,  
    include_top=True,  
    weights='imagenet',  
    input_tensor=None,  
    pooling=None,  
    classes=1000,  
    classifier_activation='softmax'  
)  
  
preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
```

```
[11]: rescale = tf.keras.layers.Rescaling(1./127.5, offset=-1)
```

```
[12]: IMG_SIZE = (100, 100)
```

```
[13]: IMG_SHAPE = IMG_SIZE + (3,)
base_model = tf.keras.applications.MobileNetV2(input_shape=IMG_SHAPE,
                                                include_top=False,
                                                weights='imagenet')
```

WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

```
[14]: image_batch, label_batch = next(iter(train_ds))
feature_batch = base_model(image_batch)
print(feature_batch.shape)
```

(32, 4, 4, 1280)

```
[15]: base_model.trainable = False
```

```
[16]: base_model.summary()
```

Model: "mobilenetv2_1.00_224"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 100, 100, 3)]	0	[]
Conv1 (Conv2D)	(None, 50, 50, 32)	864	['input_2[0][0]']
bn_Conv1 (BatchNormalization)	(None, 50, 50, 32)	128	['Conv1[0][0]']
Conv1_relu (ReLU)	(None, 50, 50, 32)	0	['bn_Conv1[0][0]']
expanded_conv_depthwise (DepthwiseConv2D)	(None, 50, 50, 32)	288	['Conv1_relu[0][0]']
expanded_conv_depthwise_BN (BatchNormalization)	(None, 50, 50, 32)	128	['expanded_conv_depthwise[0][0]']

```

expanded_conv_depthwise_relu ( (None, 50, 50, 32) 0
['expanded_conv_depthwise_BN[0][0
ReLU) ]']

expanded_conv_project (Conv2D) (None, 50, 50, 16) 512
['expanded_conv_depthwise_relu[0]

[0]']

expanded_conv_project_BN (BatchNormal (None, 50, 50, 16) 64
['expanded_conv_project[0][0]']
ization)

block_1_expand (Conv2D) (None, 50, 50, 96) 1536
['expanded_conv_project_BN[0][0]']

]

block_1_expand_BN (BatchNormal (None, 50, 50, 96) 384
['block_1_expand[0][0]']
ization)

block_1_expand_relu (ReLU) (None, 50, 50, 96) 0
['block_1_expand_BN[0][0]']

block_1_pad (ZeroPadding2D) (None, 51, 51, 96) 0
['block_1_expand_relu[0][0]']

block_1_depthwise (DepthwiseConv2D) (None, 25, 25, 96) 864
['block_1_pad[0][0]']
nv2D)

block_1_depthwise_BN (BatchNormalization) (None, 25, 25, 96) 384
['block_1_depthwise[0][0]']
malization)

block_1_depthwise_relu (ReLU) (None, 25, 25, 96) 0
['block_1_depthwise_BN[0][0]']

block_1_project (Conv2D) (None, 25, 25, 24) 2304
['block_1_depthwise_relu[0][0]']

block_1_project_BN (BatchNormalization) (None, 25, 25, 24) 96
['block_1_project[0][0]']
lization)

block_2_expand (Conv2D) (None, 25, 25, 144) 3456
['block_1_project_BN[0][0]']

block_2_expand_BN (BatchNormalization) (None, 25, 25, 144) 576

```

```

['block_2_expand[0][0]']
ization)

block_2_expand_relu (ReLU)      (None, 25, 25, 144) 0
['block_2_expand_BN[0][0]']

block_2_depthwise (DepthwiseCo (None, 25, 25, 144) 1296
['block_2_expand_relu[0][0]']
nv2D)

block_2_depthwise_BN (BatchNor (None, 25, 25, 144) 576
['block_2_depthwise[0][0]']
malization)

block_2_depthwise_relu (ReLU)   (None, 25, 25, 144) 0
['block_2_depthwise_BN[0][0]']

block_2_project (Conv2D)        (None, 25, 25, 24) 3456
['block_2_depthwise_relu[0][0]']

block_2_project_BN (BatchNorma (None, 25, 25, 24) 96
['block_2_project[0][0]']
lization)

block_2_add (Add)               (None, 25, 25, 24) 0
['block_1_project_BN[0][0]',
'block_2_project_BN[0][0]']

block_3_expand (Conv2D)         (None, 25, 25, 144) 3456
['block_2_add[0][0]']

block_3_expand_BN (BatchNormal (None, 25, 25, 144) 576
['block_3_expand[0][0]']
ization)

block_3_expand_relu (ReLU)      (None, 25, 25, 144) 0
['block_3_expand_BN[0][0]']

block_3_pad (ZeroPadding2D)     (None, 27, 27, 144) 0
['block_3_expand_relu[0][0]']

block_3_depthwise (DepthwiseCo (None, 13, 13, 144) 1296
['block_3_pad[0][0]']
nv2D)

block_3_depthwise_BN (BatchNor (None, 13, 13, 144) 576
['block_3_depthwise[0][0]']
malization)

```



```

block_3_depthwise_relu (ReLU) (None, 13, 13, 144) 0
['block_3_depthwise_BN[0][0]']

block_3_project (Conv2D) (None, 13, 13, 32) 4608
['block_3_depthwise_relu[0][0]']

block_3_project_BN (BatchNormal (None, 13, 13, 32) 128
['block_3_project[0][0]']
alization)

block_4_expand (Conv2D) (None, 13, 13, 192) 6144
['block_3_project_BN[0][0]']

block_4_expand_BN (BatchNormal (None, 13, 13, 192) 768
['block_4_expand[0][0]']
alization)

block_4_expand_relu (ReLU) (None, 13, 13, 192) 0
['block_4_expand_BN[0][0]']

block_4_depthwise (DepthwiseCo (None, 13, 13, 192) 1728
['block_4_expand_relu[0][0]']
nv2D)

block_4_depthwise_BN (BatchNor (None, 13, 13, 192) 768
['block_4_depthwise[0][0]']
malization)

block_4_depthwise_relu (ReLU) (None, 13, 13, 192) 0
['block_4_depthwise_BN[0][0]']

block_4_project (Conv2D) (None, 13, 13, 32) 6144
['block_4_depthwise_relu[0][0]']

block_4_project_BN (BatchNorma (None, 13, 13, 32) 128
['block_4_project[0][0]']
alization)

block_4_add (Add) (None, 13, 13, 32) 0
['block_3_project_BN[0][0]',
'block_4_project_BN[0][0]']

block_5_expand (Conv2D) (None, 13, 13, 192) 6144
['block_4_add[0][0]']

block_5_expand_BN (BatchNormal (None, 13, 13, 192) 768
['block_5_expand[0][0]']

```

```

ization)

block_5_expand_relu (ReLU)      (None, 13, 13, 192)  0
['block_5_expand_BN[0][0]']

block_5_depthwise (DepthwiseCo (None, 13, 13, 192) 1728
['block_5_expand_relu[0][0]']
nv2D)

block_5_depthwise_BN (BatchNor (None, 13, 13, 192) 768
['block_5_depthwise[0][0]']
malization)

block_5_depthwise_relu (ReLU)  (None, 13, 13, 192)  0
['block_5_depthwise_BN[0][0]']

block_5_project (Conv2D)       (None, 13, 13, 32)   6144
['block_5_depthwise_relu[0][0]']

block_5_project_BN (BatchNorma (None, 13, 13, 32) 128
['block_5_project[0][0]']
lization)

block_5_add (Add)              (None, 13, 13, 32)   0
['block_4_add[0][0]',
'block_5_project_BN[0][0]']

block_6_expand (Conv2D)        (None, 13, 13, 192) 6144
['block_5_add[0][0]']

block_6_expand_BN (BatchNormal (None, 13, 13, 192) 768
['block_6_expand[0][0]']
ization)

block_6_expand_relu (ReLU)     (None, 13, 13, 192)  0
['block_6_expand_BN[0][0]']

block_6_pad (ZeroPadding2D)    (None, 15, 15, 192)  0
['block_6_expand_relu[0][0]']

block_6_depthwise (DepthwiseCo (None, 7, 7, 192)   1728
['block_6_pad[0][0]']
nv2D)

block_6_depthwise_BN (BatchNor (None, 7, 7, 192)   768
['block_6_depthwise[0][0]']
malization)

```

block_6_depthwise_relu (ReLU)	(None, 7, 7, 192)	0
['block_6_depthwise_BN[0][0]']		
block_6_project (Conv2D)	(None, 7, 7, 64)	12288
['block_6_depthwise_relu[0][0]']		
block_6_project_BN (BatchNormal	(None, 7, 7, 64)	256
lization)	['block_6_project[0][0]']	
block_7_expand (Conv2D)	(None, 7, 7, 384)	24576
['block_6_project_BN[0][0]']		
block_7_expand_BN (BatchNormal	(None, 7, 7, 384)	1536
lization)	['block_7_expand[0][0]']	
block_7_expand_relu (ReLU)	(None, 7, 7, 384)	0
['block_7_expand_BN[0][0]']		
block_7_depthwise (DepthwiseCo	(None, 7, 7, 384)	3456
nv2D)	['block_7_expand_relu[0][0]']	
block_7_depthwise_BN (BatchNor	(None, 7, 7, 384)	1536
malization)	['block_7_depthwise[0][0]']	
block_7_depthwise_relu (ReLU)	(None, 7, 7, 384)	0
['block_7_depthwise_BN[0][0]']		
block_7_project (Conv2D)	(None, 7, 7, 64)	24576
['block_7_depthwise_relu[0][0]']		
block_7_project_BN (BatchNorma	(None, 7, 7, 64)	256
lization)	['block_7_project[0][0]']	
block_7_add (Add)	(None, 7, 7, 64)	0
['block_6_project_BN[0][0]', 'block_7_project_BN[0][0]']		
block_8_expand (Conv2D)	(None, 7, 7, 384)	24576
['block_7_add[0][0]']		
block_8_expand_BN (BatchNormal	(None, 7, 7, 384)	1536
lization)	['block_8_expand[0][0]']	

block_8_expand_relu (ReLU) ['block_8_expand_BN[0][0]']	(None, 7, 7, 384)	0
block_8_depthwise (DepthwiseCo nv2D) ['block_8_expand_relu[0][0]']	(None, 7, 7, 384)	3456
block_8_depthwise_BN (BatchNor malization) ['block_8_depthwise[0][0]']	(None, 7, 7, 384)	1536
block_8_depthwise_relu (ReLU) ['block_8_depthwise_BN[0][0]']	(None, 7, 7, 384)	0
block_8_project (Conv2D) ['block_8_depthwise_relu[0][0]']	(None, 7, 7, 64)	24576
block_8_project_BN (BatchNorma lization) ['block_8_project[0][0]']	(None, 7, 7, 64)	256
block_8_add (Add) ['block_7_add[0][0]', 'block_8_project_BN[0][0]']	(None, 7, 7, 64)	0
block_9_expand (Conv2D) ['block_8_add[0][0]']	(None, 7, 7, 384)	24576
block_9_expand_BN (BatchNormal ization) ['block_9_expand[0][0]']	(None, 7, 7, 384)	1536
block_9_expand_relu (ReLU) ['block_9_expand_BN[0][0]']	(None, 7, 7, 384)	0
block_9_depthwise (DepthwiseCo nv2D) ['block_9_expand_relu[0][0]']	(None, 7, 7, 384)	3456
block_9_depthwise_BN (BatchNor malization) ['block_9_depthwise[0][0]']	(None, 7, 7, 384)	1536
block_9_depthwise_relu (ReLU) ['block_9_depthwise_BN[0][0]']	(None, 7, 7, 384)	0
block_9_project (Conv2D)	(None, 7, 7, 64)	24576

```

['block_9_depthwise_relu[0][0]']

block_9_project_BN (BatchNorma (None, 7, 7, 64) 256
['block_9_project[0][0]']
lization)

block_9_add (Add) (None, 7, 7, 64) 0
['block_8_add[0][0]',
'block_9_project_BN[0][0]']

block_10_expand (Conv2D) (None, 7, 7, 384) 24576
['block_9_add[0][0]']

block_10_expand_BN (BatchNorma (None, 7, 7, 384) 1536
['block_10_expand[0][0]']
lization)

block_10_expand_relu (ReLU) (None, 7, 7, 384) 0
['block_10_expand_BN[0][0]']

block_10_depthwise (DepthwiseC (None, 7, 7, 384) 3456
['block_10_expand_relu[0][0]']
onv2D)

block_10_depthwise_BN (BatchNo (None, 7, 7, 384) 1536
['block_10_depthwise[0][0]']
rmalization)

block_10_depthwise_relu (ReLU) (None, 7, 7, 384) 0
['block_10_depthwise_BN[0][0]']

block_10_project (Conv2D) (None, 7, 7, 96) 36864
['block_10_depthwise_relu[0][0]']

block_10_project_BN (BatchNorm (None, 7, 7, 96) 384
['block_10_project[0][0]']
alization)

block_11_expand (Conv2D) (None, 7, 7, 576) 55296
['block_10_project_BN[0][0]']

block_11_expand_BN (BatchNorma (None, 7, 7, 576) 2304
['block_11_expand[0][0]']
lization)

block_11_expand_relu (ReLU) (None, 7, 7, 576) 0
['block_11_expand_BN[0][0]']

```

block_11_depthwise (DepthwiseC ['block_11_expand_relu[0][0]'] onv2D)	(None, 7, 7, 576)	5184
block_11_depthwise_BN (BatchNo ['block_11_depthwise[0][0]'] rmalization)	(None, 7, 7, 576)	2304
block_11_depthwise_relu (ReLU) ['block_11_depthwise_BN[0][0]']	(None, 7, 7, 576)	0
block_11_project (Conv2D) ['block_11_depthwise_relu[0][0]']	(None, 7, 7, 96)	55296
block_11_project_BN (BatchNorm ['block_11_project[0][0]'] alization)	(None, 7, 7, 96)	384
block_11_add (Add) ['block_10_project_BN[0][0]', 'block_11_project_BN[0][0]']	(None, 7, 7, 96)	0
block_12_expand (Conv2D) ['block_11_add[0][0]']	(None, 7, 7, 576)	55296
block_12_expand_BN (BatchNorma ['block_12_expand[0][0]'] lization)	(None, 7, 7, 576)	2304
block_12_expand_relu (ReLU) ['block_12_expand_BN[0][0]']	(None, 7, 7, 576)	0
block_12_depthwise (DepthwiseC ['block_12_expand_relu[0][0]'] onv2D)	(None, 7, 7, 576)	5184
block_12_depthwise_BN (BatchNo ['block_12_depthwise[0][0]'] rmalization)	(None, 7, 7, 576)	2304
block_12_depthwise_relu (ReLU) ['block_12_depthwise_BN[0][0]']	(None, 7, 7, 576)	0
block_12_project (Conv2D) ['block_12_depthwise_relu[0][0]']	(None, 7, 7, 96)	55296
block_12_project_BN (BatchNorm ['block_12_project[0][0]']	(None, 7, 7, 96)	384

alization)		
block_12_add (Add) ['block_11_add[0][0]', 'block_12_project_BN[0][0]']	(None, 7, 7, 96)	0
block_13_expand (Conv2D) ['block_12_add[0][0]']	(None, 7, 7, 576)	55296
block_13_expand_BN (BatchNorma ['block_13_expand[0][0]'] lization)	(None, 7, 7, 576)	2304
block_13_expand_relu (ReLU) ['block_13_expand_BN[0][0]']	(None, 7, 7, 576)	0
block_13_pad (ZeroPadding2D) ['block_13_expand_relu[0][0]']	(None, 9, 9, 576)	0
block_13_depthwise (DepthwiseC ['block_13_pad[0][0]'] onv2D)	(None, 4, 4, 576)	5184
block_13_depthwise_BN (BatchNo ['block_13_depthwise[0][0]'] rmalization)	(None, 4, 4, 576)	2304
block_13_depthwise_relu (ReLU) ['block_13_depthwise_BN[0][0]']	(None, 4, 4, 576)	0
block_13_project (Conv2D) ['block_13_depthwise_relu[0][0]']	(None, 4, 4, 160)	92160
block_13_project_BN (BatchNorm ['block_13_project[0][0]'] alization)	(None, 4, 4, 160)	640
block_14_expand (Conv2D) ['block_13_project_BN[0][0]']	(None, 4, 4, 960)	153600
block_14_expand_BN (BatchNorma ['block_14_expand[0][0]'] lization)	(None, 4, 4, 960)	3840
block_14_expand_relu (ReLU) ['block_14_expand_BN[0][0]']	(None, 4, 4, 960)	0
block_14_depthwise (DepthwiseC	(None, 4, 4, 960)	8640

```

['block_14_expand_relu[0][0]']
onv2D)

block_14_depthwise_BN (BatchNo (None, 4, 4, 960) 3840
['block_14_depthwise[0][0]']
rmalization)

block_14_depthwise_relu (ReLU) (None, 4, 4, 960) 0
['block_14_depthwise_BN[0][0]']

block_14_project (Conv2D) (None, 4, 4, 160) 153600
['block_14_depthwise_relu[0][0]']

block_14_project_BN (BatchNorm (None, 4, 4, 160) 640
['block_14_project[0][0]']
alization)

block_14_add (Add) (None, 4, 4, 160) 0
['block_13_project_BN[0][0]',
'block_14_project_BN[0][0]']

block_15_expand (Conv2D) (None, 4, 4, 960) 153600
['block_14_add[0][0]']

block_15_expand_BN (BatchNorma (None, 4, 4, 960) 3840
['block_15_expand[0][0]']
lization)

block_15_expand_relu (ReLU) (None, 4, 4, 960) 0
['block_15_expand_BN[0][0]']

block_15_depthwise (DepthwiseC (None, 4, 4, 960) 8640
['block_15_expand_relu[0][0]']
onv2D)

block_15_depthwise_BN (BatchNo (None, 4, 4, 960) 3840
['block_15_depthwise[0][0]']
rmalization)

block_15_depthwise_relu (ReLU) (None, 4, 4, 960) 0
['block_15_depthwise_BN[0][0]']

block_15_project (Conv2D) (None, 4, 4, 160) 153600
['block_15_depthwise_relu[0][0]']

block_15_project_BN (BatchNorm (None, 4, 4, 160) 640
['block_15_project[0][0]']
alization)

```


block_15_add (Add)	(None, 4, 4, 160)	0
['block_14_add[0][0]', 'block_15_project_BN[0][0]']		
block_16_expand (Conv2D)	(None, 4, 4, 960)	153600
['block_15_add[0][0]']		
block_16_expand_BN (BatchNormaliza- tion)	(None, 4, 4, 960)	3840
['block_16_expand[0][0]']		
block_16_expand_relu (ReLU)	(None, 4, 4, 960)	0
['block_16_expand_BN[0][0]']		
block_16_depthwise (DepthwiseC onv2D)	(None, 4, 4, 960)	8640
['block_16_expand_relu[0][0]']		
block_16_depthwise_BN (BatchNo rmalization)	(None, 4, 4, 960)	3840
['block_16_depthwise[0][0]']		
block_16_depthwise_relu (ReLU)	(None, 4, 4, 960)	0
['block_16_depthwise_BN[0][0]']		
block_16_project (Conv2D)	(None, 4, 4, 320)	307200
['block_16_depthwise_relu[0][0]']		
block_16_project_BN (BatchNorm alization)	(None, 4, 4, 320)	1280
['block_16_project[0][0]']		
Conv_1 (Conv2D)	(None, 4, 4, 1280)	409600
['block_16_project_BN[0][0]']		
Conv_1_bn (BatchNormalization)	(None, 4, 4, 1280)	5120
['Conv_1[0][0]']		
out_relu (ReLU)	(None, 4, 4, 1280)	0
['Conv_1_bn[0][0]']		

```

=====
Total params: 2,257,984
Trainable params: 0
Non-trainable params: 2,257,984
-----

```

```
[30]: global_average_layer = tf.keras.layers.GlobalAveragePooling2D()  
      feature_batch_average = global_average_layer(feature_batch)  
      print(feature_batch_average.shape)
```

(32, 1280)

```
[31]: prediction_layer = tf.keras.layers.Dense(3)  
      prediction_batch = prediction_layer(feature_batch_average)  
      print(prediction_batch.shape)
```

(32, 3)

```
[32]: inputs = tf.keras.Input(shape=(100, 100, 3))  
      x = data_augmentation(inputs)  
      x = preprocess_input(x)  
      x = base_model(x, training=False)  
      x = global_average_layer(x)  
      x = tf.keras.layers.Dropout(0.2)(x)  
      outputs = prediction_layer(x)  
      model = tf.keras.Model(inputs, outputs)
```

```
[41]: base_learning_rate = 0.0001  
      model.compile(optimizer=tf.keras.optimizers.  
        ↳Adam(learning_rate=base_learning_rate),  
                  loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),  
                  metrics=['accuracy'])
```

```
[42]: model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 100, 100, 3)]	0
sequential (Sequential)	(None, 100, 100, 3)	0
tf.math.truediv_1 (TFOpLamb da)	(None, 100, 100, 3)	0
tf.math.subtract_1 (TFOpLam bda)	(None, 100, 100, 3)	0
mobilenetv2_1.00_224 (Funct ional)	(None, 4, 4, 1280)	2257984

global_average_pooling2d_2	(None, 1280)	0
(GlobalAveragePooling2D)		
dropout_1 (Dropout)	(None, 1280)	0
dense_1 (Dense)	(None, 3)	3843

```

=====
Total params: 2,261,827
Trainable params: 3,843
Non-trainable params: 2,257,984
-----

```

```
[43]: len(model.trainable_variables)
```

```
[43]: 2
```

```
[44]: initial_epochs = 10

print(val_ds)

loss0, accuracy0 = model.evaluate(val_ds)
```

```

<SkipDataset element_spec=(TensorSpec(shape=(None, 100, 100, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None, 3), dtype=tf.float32,
name=None))>
8/8 [=====] - 4s 257ms/step - loss: 1.5204 - accuracy:
0.4043

```

```
[45]: print("initial loss: {:.2f}".format(loss0))
print("initial accuracy: {:.2f}".format(accuracy0))
```

```

initial loss: 1.52
initial accuracy: 0.40

```

```
[46]: history = model.fit(train_ds,
                           epochs=initial_epochs,
                           validation_data=val_ds)
```

```

Epoch 1/10
37/37 [=====] - 19s 414ms/step - loss: 0.8191 -
accuracy: 0.6723 - val_loss: 0.5107 - val_accuracy: 0.8130
Epoch 2/10
37/37 [=====] - 17s 450ms/step - loss: 0.2607 -
accuracy: 0.9075 - val_loss: 0.2060 - val_accuracy: 0.9391
Epoch 3/10
37/37 [=====] - 16s 420ms/step - loss: 0.1120 -
accuracy: 0.9652 - val_loss: 0.1159 - val_accuracy: 0.9609

```

```

Epoch 4/10
37/37 [=====] - 16s 429ms/step - loss: 0.0675 -
accuracy: 0.9856 - val_loss: 0.0660 - val_accuracy: 1.0000
Epoch 5/10
37/37 [=====] - 17s 457ms/step - loss: 0.0477 -
accuracy: 0.9932 - val_loss: 0.0447 - val_accuracy: 1.0000
Epoch 6/10
37/37 [=====] - 16s 418ms/step - loss: 0.0344 -
accuracy: 0.9941 - val_loss: 0.0314 - val_accuracy: 1.0000
Epoch 7/10
37/37 [=====] - 18s 490ms/step - loss: 0.0222 -
accuracy: 0.9992 - val_loss: 0.0241 - val_accuracy: 1.0000
Epoch 8/10
37/37 [=====] - 16s 414ms/step - loss: 0.0163 -
accuracy: 0.9992 - val_loss: 0.0192 - val_accuracy: 1.0000
Epoch 9/10
37/37 [=====] - 16s 414ms/step - loss: 0.0171 -
accuracy: 0.9975 - val_loss: 0.0171 - val_accuracy: 1.0000
Epoch 10/10
37/37 [=====] - 16s 412ms/step - loss: 0.0156 -
accuracy: 0.9975 - val_loss: 0.0117 - val_accuracy: 1.0000

```

```


[47]: acc = history.history['accuracy']
      val_acc = history.history['val_accuracy']

      loss = history.history['loss']
      val_loss = history.history['val_loss']

      plt.figure(figsize=(8, 8))
      plt.subplot(2, 1, 1)
      plt.plot(acc, label='Training Accuracy')
      plt.plot(val_acc, label='Validation Accuracy')
      plt.legend(loc='lower right')
      plt.ylabel('Accuracy')
      plt.ylim([min(plt.ylim()),1])
      plt.title('Training and Validation Accuracy')

      plt.subplot(2, 1, 2)
      plt.plot(loss, label='Training Loss')
      plt.plot(val_loss, label='Validation Loss')
      plt.legend(loc='upper right')
      plt.ylabel('Cross Entropy')
      plt.ylim([0,1])
      plt.title('Training and Validation Loss')
      plt.xlabel('epoch')
      plt.show()

```



```
[48]: base_model.trainable = True
```

```
[49]: # Let's take a look to see how many layers are in the base model
print("Number of layers in the base model: ", len(base_model.layers))

# Fine-tune from this layer onwards
fine_tune_at = 100

# Freeze all the layers before the `fine_tune_at` layer
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False
```

Number of layers in the base model: 154

```
[58]: model.compile(loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
                optimizer = tf.keras.optimizers.
                ↳RMSprop(learning_rate=base_learning_rate/10),
                metrics=['accuracy'])
```

```
[59]: model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 100, 100, 3)]	0
sequential (Sequential)	(None, 100, 100, 3)	0
tf.math.truediv_1 (TFOpLamb da)	(None, 100, 100, 3)	0
tf.math.subtract_1 (TFOpLam bda)	(None, 100, 100, 3)	0
mobilenetv2_1.00_224 (Funct ional)	(None, 4, 4, 1280)	2257984
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 1280)	0
dropout_1 (Dropout)	(None, 1280)	0
dense_1 (Dense)	(None, 3)	3843

=====
Total params: 2,261,827
Trainable params: 1,865,283
Non-trainable params: 396,544
=====

```
[60]: len(model.trainable_variables)
```

```
[60]: 56
```

```
[61]: fine_tune_epochs = 10
total_epochs = initial_epochs + fine_tune_epochs

history_fine = model.fit(train_ds,
                        epochs=total_epochs,
                        initial_epoch=history.epoch[-1],
```

```
validation_data=val_ds)
```

```
Epoch 10/20
37/37 [=====] - 58s 691ms/step - loss: 0.0025 -
accuracy: 1.0000 - val_loss: 4.2921e-04 - val_accuracy: 1.0000
Epoch 11/20
37/37 [=====] - 26s 682ms/step - loss: 1.2336e-04 -
accuracy: 1.0000 - val_loss: 1.1158e-04 - val_accuracy: 1.0000
Epoch 12/20
37/37 [=====] - 28s 746ms/step - loss: 3.3883e-05 -
accuracy: 1.0000 - val_loss: 1.5006e-05 - val_accuracy: 1.0000
Epoch 13/20
37/37 [=====] - 25s 674ms/step - loss: 5.0725e-05 -
accuracy: 1.0000 - val_loss: 3.2375e-06 - val_accuracy: 1.0000
Epoch 14/20
37/37 [=====] - 26s 682ms/step - loss: 2.5234e-06 -
accuracy: 1.0000 - val_loss: 6.1263e-07 - val_accuracy: 1.0000
Epoch 15/20
37/37 [=====] - 25s 679ms/step - loss: 1.7295e-06 -
accuracy: 1.0000 - val_loss: 2.7988e-08 - val_accuracy: 1.0000
Epoch 16/20
37/37 [=====] - 25s 676ms/step - loss: 1.0980e-07 -
accuracy: 1.0000 - val_loss: 1.8141e-08 - val_accuracy: 1.0000
Epoch 17/20
37/37 [=====] - 25s 681ms/step - loss: 4.5538e-08 -
accuracy: 1.0000 - val_loss: 1.5031e-08 - val_accuracy: 1.0000
Epoch 18/20
37/37 [=====] - 27s 736ms/step - loss: 9.0266e-08 -
accuracy: 1.0000 - val_loss: 2.0732e-09 - val_accuracy: 1.0000
Epoch 19/20
37/37 [=====] - 25s 679ms/step - loss: 1.2245e-08 -
accuracy: 1.0000 - val_loss: 5.7013e-09 - val_accuracy: 1.0000
Epoch 20/20
37/37 [=====] - 25s 677ms/step - loss: 2.0239e-09 -
accuracy: 1.0000 - val_loss: 3.6281e-09 - val_accuracy: 1.0000
```

```
[62]: acc += history_fine.history['accuracy']
      val_acc += history_fine.history['val_accuracy']

      loss += history_fine.history['loss']
      val_loss += history_fine.history['val_loss']
```


```
[64]: plt.figure(figsize=(8, 8))
      plt.subplot(2, 1, 1)
      plt.plot(acc, label='Training Accuracy')
      plt.plot(val_acc, label='Validation Accuracy')
      plt.ylim([min(plt.ylim()), 1.0])
```

```

plt.plot([initial_epochs-1,initial_epochs-1],
         plt.ylim(), label='Start Fine Tuning')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.ylim([0, 1.0])
plt.plot([initial_epochs-1,initial_epochs-1],
         plt.ylim(), label='Start Fine Tuning')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()

```




```
[65]: loss, accuracy = model.evaluate(test_dataset)
      print('Test accuracy :', accuracy)
```

```
2/2 [=====] - 1s 240ms/step - loss: 0.0000e+00 -
accuracy: 1.0000
Test accuracy : 1.0
```

4 Analysis

The best performing model was the CNN application to the sequential model - this is most likely due to the nature of the dataset itself, since it's small and with limited classes. The Pretrained Model and Transfer Learning model had a similar result but took significantly more time to process - this model is probably better for less focused datasets since it takes more time but is more accurate. They both did significantly better than the basic sequential model - which is why the sequential model has these additional parameters and options to better fine tune it towards a particular dataset.